

TTK4145 Notes

Mikkel Tiller

May 19, 2019

Part I

Code quality

Ultimate SW quality metric: **maintainability**. **Readability**, among other things, is also important.

1 Modules

Intro:

- You should be able to use modules without knowledge of its internals.
- You should be able to maintain a module without knowledge of the rest of the system.
- Composition: Easy to build supermodules from submodules, this provides scalability.
- Coupling: Weak coupling between modules.
- Cohesion: The parts of a module should be well connected.
- Interfaces should be complete and minimal.

Checklist, **encapsulation**:

1. Minimize accessibility to data members
2. Avoid exposing member data (same as above?)
3. Hide implementation details from other modules
4. Avoid making assumptions about users (other modules, derived classes)
5. Independent/loosely coupled

Checklist, other implementation issues:

1. ~ 7 data members or fewer
2. Minimize direct and indirect routine calls to other modules, i.e. collaborate with other modules only as much as absolutely necessary
3. Initialize member data appropriately
4. **Language-specific** issues should also be investigated

Key points:

1. A module interface should provide a consistent abstraction (a way to view a possibly complex implementation in a simpler form).
2. Modules are the primary tool for managing complexity.

Other tips:

- Get-functions hurt abstraction, as nothing is 'hidden' by them.
- Services often come in pairs (put/get, init/shutdown, etc. etc. etc.)
- There is no shame in having to redesign.
- Make the interface *programmatic* instead of *semantic*.

The programmatic part is the data types and other attributes that can be enforced by the compiler, while the semantic part consists of the assumptions about how the interface is supposed to be used. The semantic interface includes things like “routine A must be called before B” or “routine A will crash unless data member 1 is initialized”.

2 Routines/functions

1. Describing name
2. Does one thing, and does it well (functional cohesion)
3. Readable, easy to see that it works correctly by only reading it (and no other routines)
4. Defensive, i.e. protects itself from bad usage and inputs (assumptions are programmatic, i.e. the function can handle any input the compiler allows)
5. Less than 7 parameters

A good routine name

1. Describes 'all' the routine does

2. Is specific
3. Describes the return value, if it is a function
4. Consists of a verb and an object, if it is a procedure
5. Hot tip: Standardize conventions for names and abbreviations, as well as parameter order

Reasons to create a routine:

- Encapsulate something
- Introduce abstraction
- Improve readability
- Avoid duplicating code
- Simplify testing

Big-picture issues

- Loose coupling with other routines.
- Length determined by functionality and logic, not some artificial coding standard.
- The parameter list provides a consistent interface abstraction.

Key points

- Routines should improve readability, reliability and modifiability.
- Functionally cohesive (does one thing, and does it well).
- The name is a metric of quality, a good name expresses what the routine does.
- Macro routines are a last resort.

3 Variables

A good variable name is

- Describing and specific
- Adapted to the scope and data type of the variable

Remember: Good code is self-documenting. If things are unclear after reading the code, focus on improving the code, rather than adding comments! Comments should express intent or summarize, not repeat the code. Comments add to the size of the project, remember that they also add to maintenance cost.

Part II

Fault tolerance

Four sources of faults in embedded systems:

1. Inadequate specification, i.e. misunderstanding the interactions between the program and environment.
2. Design errors in software components. We typically don't know the consequences of these.
3. Failure of hardware components. More predictable than the one above.
4. Interference in the supporting communication subsystem.

4 Reliability, failure and faults

Reliability is a measure of how well a system conforms to the specification of its behavior. *Response times* are an important part of the specification. **Failure** is when a system deviates from the specification. Highly reliable \iff low failure rate.

The above definitions are concerned with the systems behaviour (external appearance). Failure stems from internal errors, whose algorithmic or mechanical causes are called **faults**. This motivates the following definition of a faulty component: A faulty component is one that under certain circumstances, during the lifetime of the system, results in an error.

An external state that is not in the specification is regarded as a failure, and an internal state that is not specified is called an error. A fault is **active** when it produces an error, otherwise it is **dormant**. The error propagates through the system, and manifests itself as part of the external behavior.

A failure in one (sub)system can cause faults in other systems, as the following chain of events illustrates:

Fault $\xrightarrow{\text{activation}}$ Error $\xrightarrow{\text{propagation}}$ Failure $\xrightarrow{\text{causation}}$ Fault

With regard to time, there are three kinds of faults:

1. **Transient faults** occur at some point in time, then disappear at some later point in time. E.g. HW response to external electromagnetic field, fault disappears when the field disappears. Common kind of fault in communication systems.
2. **Permanent faults** start at a particular time, then remain in the system until they are fixed. E.g. SW design error.

3. **Intermittent faults** are transient faults that occur from time to time, common with e.g. heat sensitive HW.

Bugs are software faults, and originally there were two kinds:

- **Bohrbugs** are reproducible and identifiable, can be removed during testing or with design diversity techniques.
- **Heisenbugs** are only active under certain rare circumstances, and often disappear when investigated. An example is code shared between concurrent tasks that is not properly synchronized. Heisenbugs can result from software ageing, e.g. not freeing allocated memory, and exhausting the available memory after a long time. Restarting the system clears this bug.

5 Failure modes

A system provides services, and its failure modes can be classified according to their impact on these services. Two general classes:

- **Value failure** - Error in the value associated with service. Can be the result of a data conversion, e.g. 64-bit to 8-bit.
- **Time failure** - Service not delivered at the correct time.

Combinations of the two are called *arbitrary* failures. Time failures can result in the service being delivered:

- too early
- too late (performance error)
- infinitely late (omission failure)

How can a system fail?

1. Fail uncontrolled - arbitrary errors
2. Fail late - delivers service too late in the time domain
3. Fail silent - omission failure
4. Fail stop - same as above, but permits other systems to detect that it has failed silently
5. Fail controlled - fails in a specified controlled manner
6. Fail never - self explanatory

6 Fault prevention and fault tolerance

Fault prevention aims to eliminate any and all faults before the system goes into operation, whilst **fault tolerance** enables the system to continue functioning even in the presence of faults. Both approaches attempt to give the system well-defined failure modes.

6.1 Fault prevention

Two steps; fault avoidance and fault removal. Fault avoidance consists of:

- Acquiring reliable hardware and protecting it against interference.
- Rigorous specification of requirements, design based on e.g. UML (avsky).
- Use languages with data abstraction and modularity, like Ada and Java.

Fault removal consists of removing causes of errors, mainly by systems tests. However, testing can never remove all potential faults:

- A test can only show the presence of faults, not their absence.
- It might be impossible to test under realistic conditions.
- Errors from the requirement analysis might not become visible until the system is in operation.

Any system will fail eventually (HW or SW), and for real-time systems we need fault tolerance.

6.2 Fault tolerance

Three levels of fault tolerance:

1. **Full fault tolerance** - operation continues without significant loss of performance, but only for a limited time.
2. **Fail soft** - operation continues, but a partial degradation in performance is accepted until recovery or repair.
3. **Fail safe** - the system's integrity is maintained, but a temporary halt in operation is accepted. I.e. the system is shut down in a safe state.

Back in the day fault-tolerant design was based on three assumptions:

1. Algorithms are correctly designed.
2. All failure modes are known.

3. All possible interactions with the environment are known.

This is not realistic today, with multi-core processors and such, hence both anticipated and unanticipated errors must be accounted for.

6.3 Redundancy

Protective redundancy introduces components that detects and recovers the systems from faults, but are unnecessary for normal operation. When designing a fault-tolerant system, the goal is to minimize redundancy while maximizing reliability, subject to constraints on cost, size and power consumption. The redundant components can (and will) increase complexity, and it is useful to separate them from the rest of the system.

We separate between static and dynamic redundancy, both for hardware and software. First, let's have a look at hardware redundancy:

- **Static redundancy** (or masking) is based on redundant components “hiding” faults. An example is Triple Modular Redundancy (TMR), where a majority voting circuit is used. The output of three identical components are compared, and if one differs from the others, its output is masked out. It is assumed that faults are transient.
- **Dynamic redundancy** is an error-detection facility within a component, making it possible for that component to indicate if its output is in error. Note that the component does not hide or fix the error, that must be done by some other part of the system. Examples are check-sums (see parity byte section on Wikipedia for very simple example) and parity bits.

For fault tolerance with regards to software design, we have *N-version programming* which works like masking, and *error detection and recovery*. The latter is dynamic redundancy in the sense that recovery is only brought into action once an error has occurred.

7 N-version programming

From one initial specification, N independent programs are created. In operation, they run concurrently, and their outputs are compared by a driver process. The “correct” result is determined by majority of vote, like with masking. Challenges concerning N-version programming are:

- **Initial specification** - It is close to impossible to produce unambiguous specifications.
- A complex part of the specification can potentially induce faults for all the N independent developer teams.

- Budget concerns, N-versions are N times more expensive than one.
- **Granularity** - How often results are compared affects overhead in one direction, and fault tolerance in the other.
- **Inexact voting** - Results may not agree exactly, even when no fault has occurred. This is a consequence of *finite-precision* arithmetic and the possibility of multiple correct solutions (a quadratic eq. being a simple example). One solution is to regard values inside a range of Θ to be equal, but then the problem arises once again for values close to the boundaries of the range.

8 Software dynamic redundancy

Statically redundant components operate whether or not an error has occurred. With dynamic redundancy, however, the redundant components are only put into play when an error occurs. There are four phases to dynamic redundancy in software:

1. **Error detection**
2. **Error diagnosis** - There is a delay between a fault becoming active and error detection, the propagation of erroneous information in the system is assessed.
3. **Error recovery** - Transform the corrupted system into a state where it can continue operation.
4. **Fault treatment** - Maintenance must be performed to correct the underlying fault responsible for the error.

8.1 Error detection

There are two classes:

- **Environmental detection** - Detection by hardware (e.g. overflow error) or run-time support system (e.g. out of bounds error for array).
- **Application detection**
 - **Replication checks** - Check if results are equal for duplicate threads or processors (like N-version programming).
 - **Timing checks** - Can be a *watchdog timer* that has to be reset with a given frequency, or detection of missed deadlines by the scheduling system.
 - **Reversal checks** - Compute the input from the output, and compare with the actual input.

- **Coding checks** - E.g. checksum
- **Reasonableness checks** - `assert()`-function
- **Structural checks** - E.g. count number of elements in list to confirm integrity.
- **Dynamic reasonableness checks** - Error assumed if new output is too different from previous value.

8.2 Error diagnosis

Software designers aim to minimize the damage caused by a faulty component, this is called *firewalling*. Two techniques are:

- **Modular decomposition** - Modules only communicate through well-defined interfaces, internal details are hidden. Provides a static structure.
- **Atomic actions** are indivisible, and appear to happen instantaneously for the rest of the system. Often called **transactions** or atomic transactions. They are used to move the system from one consistent state to another, and limit the flow of information between components/modules.
- There are also **protection mechanisms** which may stop a process from accessing a resource based on its access permissions.

8.3 Error recovery

There are two approaches; **forward** and **backward** error recovery.

Forward error recovery tries to continue from an erroneous state by finding a new consistent (but probably sub-optimal) state. An example is Hamming codes (haven't read about them). Useful if the error related to the previous state may happen many times in a row, and one cannot afford to return to that state.

Backward error recovery restores the system to a previous safe state, a **recovery point**, and executes a different code block than the one that lead to an error. The new code should have the same functionality, but use a different algorithm. Setting up a recovery point is called **checkpointing**.

Backward recovery can be used to recover from unanticipated faults (very good), but cannot undo effects the fault had on the environment (e.g. launching a missile). Furthermore, it may be costly in a real-time sense to save state from time-varying sensor data.

State restoration with concurrent processes is not necessarily simple, as is seen from the **domino effect**. Let us say we have two processes, say P_1 and P_2 that communicate, synchronize and set up recovery points. If P_1 detects an error at time T_e , it will roll back to its previous safe state. But what if there was communication between P_1 and P_2 between that state and

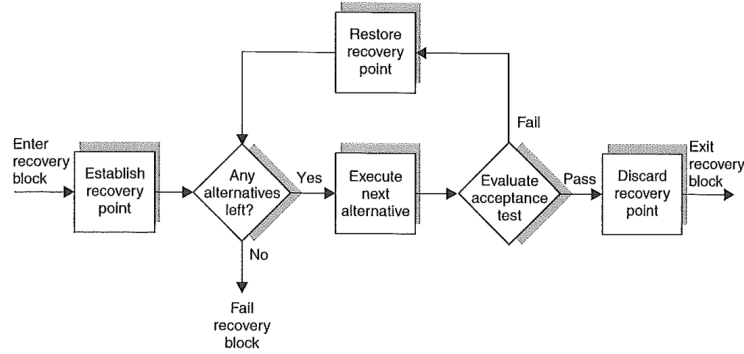


Figure 1: Structure of a recovery block.

T_e ? Then P_2 must also roll back to its previous state, and there might have been some communication in between there as well. This might continue until both processes are back to square one, and is called the *domino effect*.

The probability of the domino effect increases with the number of concurrent processes. A **recovery line**, a consistent set of recovery points, is required to avoid the effect.

9 Recovery blocks

Figure 1 illustrates a recovery block.

9.1 Acceptance test

The acceptance test provides an error detection mechanism, e.g. with invocations of `assert()`. There is always a trade-off between a comprehensive test and affecting the 'happy path' as little as possible. All the error detection techniques discussed in Section 8.1 can be used to create acceptance tests.

Acceptance tests do not focus on specific error situations and error returns, since these will never handle unexpected errors. Instead we put demands on the current state to be acceptable.

10 Comparison between N-version programming and recovery blocks

Brief comparison:

- N-version is static, all versions run regardless of whether an error has occurred or not. Recovery blocks are dynamic.

- N-version requires a driver process, while recovery blocks need an acceptance test. At run-time N-version requires N times the resources, since recovery blocks only run one code block at a time. However, establishing recovery points, and reverting to them is expensive.
- Both are prone to errors stemming from ambiguous specifications.
- Acceptance tests may be more flexible than e.g. an inexact voting scheme.
- The backward error recovery of recovery blocks cannot undo effects on the environment (not impossible to design a system that avoids this problem), whereas N-version requires everything to go through the driver before it can affect the outside world.

11 Dynamic redundancy and exceptions

Exception := the occurrence of an error. Telling the 'invoker' about the error condition is called **raising** (or signaling or throwing) an exception, and the invoker then **handles** (or catches) the exception.

Exceptions can be regarded as forward error recovery, as the state is not rolled back, but control is handed over to the exception handler (still, backward recovery can be implemented with exceptions). There is an example in B&W that illustrates some of the controversy surrounding exceptions.

Part III

Fault model & software fault masking

Keywords from the part on basic fault tolerance are:

- Acceptance tests
- Merging failure modes
- Redundancy

For a fault tolerant system we want the following (progressively better):

1. Failfast (errors are detected immediately)
2. Reliable (failfast + the system is repaired)
3. Available (continuous operation)

An overview of the process is:

1. Find failure modes.
2. Detect errors / simplify error model, inject errors for testing
3. Error handling by redundancy \rightarrow reliable and available module.

On error injection:

1. Simplify by merging failure modes.
2. Inject failed acceptance tests.

Examples for fault tolerant modules, three basic modules: storage, communication and processing.

12 Storage

Imagine an array of data. Assume unreliable read and write functions.

12.1 Failure modes

For writing we have:

- Writes the wrong data
- Writes to the wrong place

- Does not write
- Fails

Likewise, for reading:

- Give wrong data
- Give old data
- Give data from wrong place
- Fails

12.2 Detection, merging of error modes and error injection

- Detect by also writing address, checksum, versionID and statusbit to the buffer.
- Merging: All errors → Fail
- For error injection, spawn a thread that runs in parallel and flips status bits. This seems really smart!

12.3 Handling with redundancy

- Keep several copies of the buffer, the one with the newest versionID is used (returned).
- Always write back when a reading error occurs (write a 'safe state'?).
-

13 Messages

13.1 Failure modes

- Lost
- Delayed
- Corrupted
- Duplicated
- Wrong recipient

13.2 Detection, merging of error modes

- Session ID
- Checksum
- Ack (acknowledgement)
- Sequence numbers
- All errors → Lost message

13.3 Handling with redundancy

- Timeout and retransmission

14 Processes/calculations

Error mode: Does not yield the next correct 'side effect'.

Detect and merge: All failed acceptance tests → STOP (failfast).

Three ways to handle with redundancy, described in the following subsections.

14.1 Checkpoint-restart

- Write state to storage after each (successful) acceptance test, before each side effect.
- This yields error containment, but requires good acceptance tests.

14.2 Process pairs

- Two processes; primary and backup (primary does the work).
- Backup takes over if (when) primary fails, and becomes new primary.
- Primary sends IAmAlive-messages and checkpoints to backup.
- NB! Because of resending this does not rely on safe communication. I.e. we get redundancy by resending, communication errors are masked out.

14.3 Persistent processes

- Transactional infrastructure
- All calculations are transactions, i.e. atomic transformations from one consistent state to another.

- The processes are then 'stateless', all states are stored in a database.
- For such simple processes OS can take care of restart.

Note that reliable and available storage, communication and calculations are necessary to make this transactional infrastructure. Hence one often falls back to one of the first two options.

Part IV

Transaction fundamentals

Context and motivation for introducing transactions:

- We need error handling and -containment for systems with multiple participants (threads, processes, distributed systems). These participants must often cooperate in the error handling.
- Transactions (and atomic actions) are techniques/frameworks that provide the means to do this. They fall under the category of dynamic SW redundancy.
- They contribute towards the desired 'error assessment and confinement' design, and help avoiding the 'domino effect'.
- A motivating example is a process control plant with many local controllers, supervisory tasks, monitoring, optimization, several modes of operation and high demands for safety.

From the learning goals, **eight design patterns**:

1. Locking
2. Two-phase commits
3. Transaction manager (TM)
4. Resource manager (RM)
5. Log
6. Checkpoints
7. Log manager
8. Lock manager

And some additional terms:

- Optimistic concurrency control
- Two-phase commit optimization
- Heuristic transactions
- Interposition

Atomic operation: Indivisible operation, either it happens or it does not at all.


(Atomic) transaction:

- All-or-nothing property to work conducted within its scope.
- Shared resources are protected.

A transaction is an atomic action with backward error recovery. ACID properties of transactions:

1. Atomicity: The transaction either commits successfully or rolls back (aborts) completely at fail.
2. Consistency: Preserve consistent state.
3. Isolation: Intermediate states during a transaction are not visible to the outside. Further, transactions appear to be executing *serially*, even when they are not.
4. Durability: The effects of a committed transaction are never lost, i.e. they are stored in stable storage, such as disk.

Two-phase commits: Associated with each transaction is a coordinator C (could also be a Transaction manager (TM)), that communicates with the participants. The message flow is as follows:



The first phase is the commit request (or voting) phase, and includes lasts until the **VOTE YES/NO** arrow in the above diagram. The second phase is the commit phase, where the operation of each participant is either completed (if all votes were yes) or rolled back (if anyone voted no). Refer to https://en.wikipedia.org/wiki/Two-phase_commit_protocol for more on the algortihm.

A disadvantage with two-phase commits is that the protocol is blocking. Participants will block after they have voted, awaiting a commit or rollback message. If the coordinator fails, they will never receive either.

15 The 8 design patterns

15.1 Locking

- Ensures that intermediate states are not propagated
- No communication outside participants of the transaction.

With backwards error recovery, the flow of control of a single-threaded program is as follows:

```
allocate locks
create recovery point
do work (goto end if any problems)
label end:
if (error) {
    roll back to recovery point
    status = FAIL
} else {
    status = OK
}
release locks
return status
```

Note that it is the **acceptance test** that is responsible for detecting the errors in the program above. The **lock manager** can:

- Release all locks associated with transaction X
- Tidy up after restart
- Handle resources shared by several RMs
- Include deadlock avoidance and detection algorithms

15.2 The transaction manager

Useful if there are multiple participants.

- Creates the transaction, keeps track of participants
- Plays the role of the coordinator discussed earlier
- Can give the transaction a deadline, and abort if it is missed

15.3 The resource manager

Resource manager \equiv transaction participant.

- Keeps track of locks and recovery points
- Communicates with manager during voting and commit phases

15.4 The log

- Every participant (including the TM) writes every planned change of state to the log, and waits until the operation is confirmed safe to perform it
- Processes can restore state after restart by executing the logrecords (of committed work)
- Logrecords can be executed backwards to undo actions, we don't need recovery points anymore!
- Solution: Write checkpoints (state and list of active transactions) to the log. Log older than the last checkpoint with no active transactions can safely be deleted
- The **log manager** queues logrecords, optimizes disk access and sends receipts (ack) on received log record

15.5 Other terms

Optimistic concurrency control:

- Assume no task-interference, detect afterwards and abort if it has happened
- Replaces locking

Two-phase commit optimizations:

- Early abort using e.g. ATC
- One-phase when there is only one participant (can commit or block directly, without voting)
- Read-only operations will never be aborted, thus it will always be a commit
- Last resource commit: One participant (maybe responsible for real-world interaction) gets to wait until everyone else has voted to make his decision

Heuristic transactions:

- What if we give our vote, then lose connection?
- We have to make a local guess, and maybe do forward error recovery after.

Interposition is the TM's ability to play the role of RM.

Part V

Atomic actions

Definition: An atomic (indivisible, isolated) transformation of the system from one consistent state to another. Transactions are generalized to atomic actions with forward error recovery. I.e. atomic actions are actions with either forward or backward error recovery, while transactions are actions with backward error recovery. Nice, because backward ER might not be viable in a RT setting.

16 Features of an AA

- Explicit membership
- Well defined **start boundary** with recovery points in the case of backward ER (not necessarily coordinated in time)
- **Side boundary**, i.e. no communication with non-participants
- A clear **end boundary** (coordinated in time!)

16.1 Standard AA implementation

Start boundary:

1. Dynamic (call to e.g. `startWork()`)

Side boundary:

1. Locking
2. No unlocking before safe state (end boundary)

End boundary:

1. Vote-counting (two-phase commit)
2. Synchronization primitive *barrier* (no can pass through before everyone has arrived)

17 Language-specific implementation

Java:

- Synchronized methods, wait, notify/notifyAll
- Asynchronous exceptions

Ada:

- Protected objects
- Functions procedures and entries with guards

C/POSIX:

- `pthread_cancel`
- `setjmp` and `longjmp`

Part VI

Shared variable synchronization

18 Semaphores

A **semaphore** is like an integer, with three key differences:

1. It can be initialized to any value, but after that only incremented or decremented (its value cannot be read).
2. When the semaphore is decremented by a thread, and the result is negative, the thread blocks.
3. When the semaphore is incremented by a thread, one waiting thread gets unblocked (if any).

Things to note:

- After one thread increments the semaphore, and another is woken, they run concurrently.
- A positive value represents the number of threads that can decrement without blocking.
- A negative number represents the number threads that have blocked and are waiting.

The basic syntax used is:

- `sem = Semaphore(1)` to create a new semaphore with the given initial value.
- `sem.signal()` to increment the semaphore (and wake a waiting thread).
- `sem.wait()` to decrement the semaphore (and block if the result is negative).

In the classical example of incrementing/decrementing `int i`, `i`'s value must be set pretty high to see synchronization errors. The reason for this is that context switching (switching between threads and saving states) does not happen until a certain time passes. For small `i`, one thread will typically finish before the second gets a chance to start.

19 Standard problems

A (data) **race condition** is a fault in the design of the interaction between two threads, which leads to the result being highly dependent on the sequence and timing of access to shared variables.

A **livelock** occurs when a thread gets stuck in e.g. a busy-waiting loop (more generally; a subset of the possible states), and is unable to proceed.

A **deadlock** occurs when the system is stuck in a circular wait, with no threads being able to proceed.

Starvation happens when a thread is unable to gain access to a resource it requires, because other threads keep getting it first. Can be caused by an 'unfair' scheduler.

A thread that is free from deadlocks, livelocks and infinite postponements (starvation) is said to possess **liveness**.

20 Monitors

What is bad with semaphores?

- Forgetting a wait can lead to multiple threads running concurrently in a **critical region** (a section of code that should always be executed under mutual exclusion).
- Forgetting a signal can lead to a deadlock.
- The code is distributed all over the system, making maintenance hard.

20.1 Critical regions

- Critical regions but the responsibility for mutual exclusion on the compiler, instead of the programmer. Critical regions of the same name mutually exclude.
- In **conditional critical regions**, the thread must wait for the mutex lock *and* check a boolean condition before it can proceed. A **guard** is a more sophisticated form of this boolean condition, which instead runs a test (procedure) that opens or closes for execution.

Limitations

- (Conditional) critical regions are still distributed throughout the program.

20.2 Monitors

- A collection of local variables and procedures (basically a module), with a mutex that only allows one thread at a time to access its methods and variables.
- Pro: All code that accesses the shared data is localized.
- Has **condition variables** (queues of threads waiting for some condition to be true) with operations **suspend** and **resume** to block and unblock threads.
- **suspend** releases the monitor lock.
- The operations are safe, because they can only be accessed/called from inside the monitor.

Cons

- Suspending in a nested call does not release the outer lock. Ex. a procedure in monitor A calls a procedure in monitor B, which contains a **suspend** statement. This makes it infeasible to build supermodules from submodules (that are monitors).
- The possibility for deadlocks and data races still exists.
- Does not provide other synchronization than condition variables.

21 Synchronization mechanisms in POSIX, Java and Ada

21.1 Java

The `synchronized` keyword is used to make methods thread-safe. Consider the following class:

```
public class MyClass {
    private int i;

    public MyClass(int initValue) {
        i = initValue;
    }

    public synchronized void increment() {
        i++;
    }
}
```


When thread A is executing `increment()`, all other threads that (wish to) invoke synchronized methods for the same `MyClass`-object block (suspend) until thread A releases the monitor lock.

`wait()` is used to suspend the current thread, like this:

```
public synchronized void conditionalIncrement() {
    while(i < 3) wait();
    i++;
}
```

Note that since the method is `synchronized`, thread A must hold the monitor lock before it can invoke it. When it calls `wait()`, it releases the lock and suspends execution. It will be woken up at some time after the following procedure has run:

```
public synchronized void importantChange() {
    i = 3;
    notifyAll();
}
```

All threads waiting on the lock are notified when `notifyAll()` is called.

21.2 Ada

`Protected types` are data types protected from inappropriate simultaneous access by multiple tasks. They are split into an interface definition and an operation implementation, as shown below.

```
protected type Counting_Semaphore is
    procedure Release;
    entry Acquire;
    function Lock_Count return Natural;
private
    Count : Natural := 0; -- Natural number, i.e. geq 0.
end Counting_Semaphore;

protected body Counting_Semaphore is
    procedure Release is
    begin
        if Count > 0 then
            Count := Count - 1;
        end if;
    end Release;

    entry Acquire when Count < 11 is
    begin
```

```

        Count := Count + 1;
    end Acquire;

    function Lock_Count return Natural is
    begin
        return Count;
    end Lock_Count;
end Counting_Semaphore;

```

We take note of the following from the example:

- **Procedures** can modify the state, and guarantee mutual exclusion.
- **Entries** are similar, but they have a boundary condition that must be true before they can execute.
- **Functions** cannot change the state, only read it. They have concurrent access, not exclusive.

Ada supports ATC, which is useful for forward error recovery. This is implemented with the **select-then-abort** statement. Tasks are notified asynchronously and immediately when an event occurs, either a **delay** or the condition of an **entry** becoming true. This allows other tasks to be aborted when one fails, so that no wasted work is done.

```

select
    Wait_For_Abort; -- Run this entry when the boundary condition is true
    -- Forward ER code
then abort
begin
    Do_Work;
exception
    when Work_Failed => -- Can be raised by Do_Work
        Abort_Procedure; -- Triggers Wait_For_Abort
end;
end select;

```

The **requeue** keyword is used in an entry body to delay the execution of certain statements in the body. Example:

```

entry My_Entry when My_Cond = True is
begin
    if Data_Available then
        Do_Work;
    else
        requeue Wait_For_Data;
    end if;
end My_Entry;

```

21.3 POSIX/C

Example usage of `pthread` with mutexes:

```
int i = 0;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;

void* increment() {
    for (int j = 0; j < 1000000; j++) {
        pthread_mutex_lock(&mutex);
        i++;
        pthread_mutex_unlock(&mutex);
    }
    return NULL;
}

void* decrement() {
    for (int j = 0; j < 1000000; j++) {
        pthread_mutex_lock(&mutex);
        i--;
        pthread_mutex_unlock(&mutex);
    }
    return NULL;
}

int main() {
    pthread_t A;
    pthread_t B;

    pthread_create(&A, NULL, increment, NULL);
    pthread_create(&B, NULL, decrement, NULL);

    pthread_join(A, NULL);
    pthread_join(B, NULL);

    printf("i = %d\n", i);
    return 0;
}
```

For exceptions and error handling in C, `goto` and `setjmp/longjmp` can be used. Basic `goto` syntax:

```
goto label;  
.  
.  
.  
label: statements // Could also be above goto
```

See <https://en.wikipedia.org/wiki/Setjmp.h> for examples with `setjmp/longjmp`.

Part VII

Modeling of concurrent programs

Classic deadlock example:

```
T1:
  while (1) {
    Wait(A)
    Wait(B)
    ...
    Signal(A)
    Signal(B)
  }
```

```
T2:
  while (1) {
    Wait(B)
    Wait(A)
    ...
    Signal(B)
    Signal(A)
  }
```

Let us introduce **processes** and **events**. All processes that 'care about' an event experience it at the same time. For modeling we use FSP (Finite State Processes) \subset CSP (Communicating Sequential Processes). Note that Ada **entries** can be modelled as events. For messagebased systems (message passing, synchronous communication), modeling with processes, events and guards comes naturally.

```
T1 = (t1wa -> t1wb -> t1sb -> t1sa -> T1).
T2 = (t2wb -> t2wa -> t2sa -> t2sb -> T2).
```

No deadlock is visible here, we need to model the semaphores:

```
SA = (t1wa -> t1sa -> SA | t2wa -> t2sa -> SA).
SB = (t1wb -> t1sb -> SB | t2wb -> t2sb -> SB).
```

To model the complete, concurrent/parallel system:

```
||SYSTEM = (T1 || T2 || SA || SB).
```

In a **transition diagram**, a deadlock is a state with no exit. A livelock is a subset of states we cannot exit. Dining philosophers with deadlock is shown next

```

FORK = (get -> put -> FORK).
PHIL = (sitdown -> right.get -> left.get -> eat ->
        right.put -> left.put -> arise -> PHIL).

|| DINERS(N=5) = forall [i:0..N-1]
    (phil[i]:PHIL ||
     {phil[i].left, phil[(i-1+N)%N].right}::FORK).

```

Some advance syntax is used here:

- `phil[i]:PHIL` exchanges `eat` with `phil1.eat`, and thus takes care of all events for all N philosophers.
- `phil[i].left, phil[(i-1+N)%N].right` means that either of the two options in curly braces are possible.
- `phil[i].left, phil[(i-1+N)%N].right::FORK` then means that either `phil[i].left.get` or `phil[(i-1+N) mod N].right.get` are possible for the same fork, i.e. it can be picked up by the philosopher to its right or left.

See notebook for transition diagram with three philisophers.

Part VIII

Scheduling

Concurrent programs exhibit non-determinism, and a real-time system needs to restrict this non-determinism through scheduling. A scheduling scheme provides two key features:

1. An algorithm for ordering the use of system (CPU) resources.
2. A way to predict the worst-case behavior of the system when the scheduling algorithm is applied.

There are **static** (prediction only before execution) and **dynamic** (run-time decisions are used) schemes. Focus for us: static schemes, specifically pre-emptive priority-based schemes on a single-processor system. The task with the highest priority will always run, unless it is suspended or delayed for some reason. The two features above can be re-stated as:

1. Priority assignment algorithm
2. Schedulability test

21.4 Non-preemptive scheduling (cyclic executive)

What is non-preemptive scheduling?

- Fixed set of tasks with fixed periods.
- Consists of a table of procedure calls, the major cycle, comprised of smaller minor cycles with fixed duration.
- At run-time, no tasks run concurrently, therefore mutual exclusion is guaranteed (we don't need e.g. semaphores).
- All task periods must be a multiple of the minor cycle period (obvious drawback).

What are some of its drawbacks?

- Difficult to incorporate sporadic tasks.
- Difficult to construct.
- Difficult to incorporate tasks with long periods
- 'Large' tasks will need to be split up into several procedures, which may hurt the quality of the code, and make it more error-prone.

21.5 Task-based scheduling

The simple task model makes the following assumptions:

1. Fixed set of periodic tasks, with known periods
2. Independent tasks
3. Deadline = period
4. Zero overhead
5. Fixed worst-case execution time
6. No internal suspension (e.g. blocking I/O request)
7. One core (CPU)

In real-time systems priority is determined by temporal requirements. Two important approaches to scheduling are:

- Fixed-priority scheduling (FPS), static
- Earliest deadline first (EDF), dynamic

21.6 Response time analysis

Response time as function of computation time and interference:

$$R_i = C_i + I_i,$$

where

$$I_i = \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil$$

Note that the equation has R_i on both sides, and is difficult to solve. We have a better way:

$$w_i^{n+1} = C_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n}{T_j} \right\rceil C_j, \quad w_i^0 = C_i$$

When $w_i^{n+1} = w_i^n$, we have the solution for R_i . If $R_i \leq T_i = D_i$, all is good.

21.7 Priority ceiling protocol

A priority ceiling protocol avoids **unbounded priority inversion** and deadlock. Each shared resource is assigned a ceiling priority equal to the highest priority of any task that can lock it. We have two variants: OCCP (original) and ICCP (immediate). Both work by temporarily raising the priorities of tasks.

In **ICCP**, a task's priority is set to the ceiling priority of the resource it locks, so that no task that may want to lock the resource can be scheduled.

In **OCCP** the priority of a low-priority task X which locks resource Q is raised (to the Q's ceiling) when a higher-priority task Y tries to acquire Q. I.e. no priority reassignment occurs before any actual blocking happens.

Note that in both cases a task can only lock a resource if its *dynamic* priority is higher than the ceilings of all resources already locked by other tasks.